

Model Results

Evan Schwartz, Ethan Wood, Max Van Fleet

2024-11-14

Loading data and Libraries

```
library(readr)
library(caret)
```

```
## Loading required package: ggplot2
## Loading required package: lattice
```

```
library(leaps)
library(glmnet)
```

```
## Loading required package: Matrix
## Loaded glmnet 4.1-8
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
## The following object is masked from 'package:dplyr':
##
##   select
```

```
library(ggplot2)
library(lattice)
```

```
data = read_csv("/Users/evanschwartz/664Project/Data/final_data_tranform_clean.csv", show_col_types = F)
```

```
## New names:
## * `` -> `...1`
```

```
#Last Minute data cleaning, some economic vairables removed due to high collinearity
data = dplyr::select(data, -c(1,4,6, 8,9, 10, 12, 13))
```

K-folds for data=Here we prepare the training data by shuffling and preforming a 90%-10% training-test split.

```
# 90/10 data
set.seed(664)

shuffleddata = data[sample(nrow(data)), ]
data_naive_shuff <- dplyr::select(shuffleddata, c(1, 2, 3, 6, 7, 10, 28:59, 70:77,))

# Calculate the index for the 90-10 split
index90 = floor(0.9 * nrow(shuffleddata))

train = shuffleddata[1:index90,]
train_naive = data_naive_shuff[1:index90,]
# Create the 10% testing data
test = shuffleddata[(index90 + 1):nrow(shuffleddata),]
test_naive = data_naive_shuff[(index90 + 1):nrow(data_naive_shuff),]

#of the 90 data, 7 fold cross validation

folds = createFolds(train$loglatestPrice, k = 7, list = TRUE)
```

AIC and BIC

```
MSE.bic = vector()
bic.num = vector()
MSE.aic = vector()
aic.num = vector()

for (j in 1:7) {
  #setting respective fold to either training or testing
  train_data = train[-folds[[j]], ] #eliminate columns in testing data
  test_data = train[folds[[j]], ]

  #doing best subset selection on our data
  regfit.train = regsubsets(loglatestPrice ~ ., data = train_data, method= "backward" )
  reg.summary = summary(regfit.train)

  #selecting which subset is the best, specifically selecting how many variables it has
  bic.number = which.min(reg.summary$bic)
  bic.num[j] = bic.number

  #creating the model with the respective number of coefficients
  bic.mdl = coef(regfit.train, bic.number)

  # testing the models that are created
  test.mat2 = model.matrix(loglatestPrice ~ ., data = test_data)
  bic.pred = test.mat2[, names(bic.mdl)] %*% bic.mdl

  # Start with the full model
```

```

full.model = lm(loglatestPrice ~ ., data = train_data)

# Perform backward stepwise selection using AIC
#stepwise.model = stepAIC(full.model, direction = "backward", trace = FALSE)

#predict with the
#aic.pred <- predict(stepwise.model, newdata = test_data)
#aic.num[j] <- length(coef(stepwise.model))

#computing the MSE for each model
MSE.bic[j] = mean((test_data$loglatestPrice - bic.pred)**2)
#MSE.aic[j] = mean((test_data$loglatestPrice - aic.pred)**2)
}

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax, force.in =
## force.in, : 1 linear dependencies found

## Reordering variables and trying again:

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax, force.in =
## force.in, : 2 linear dependencies found

## Reordering variables and trying again:

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax, force.in =
## force.in, : 2 linear dependencies found

## Reordering variables and trying again:

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax, force.in =
## force.in, : 1 linear dependencies found

## Reordering variables and trying again:
paste("Estimated MSE for BIC is ", sum(MSE.bic))

## [1] "Estimated MSE for BIC is 0.835780797232455"
#paste("Estimated MSE for AIC is ", sum(MSE.aic))

bic.num

## [1] 9 9 8 9 8 9 8
BIC says model with ____ variables, AIC says model with ____ variables

```

LASSO & Ridge

LASSO

```

set.seed(1234)

vectorMSElasso = vector()
for (j in 1:7) {

  train_data = train[-folds[[j]], ]
  test_data = train[folds[[j]], ]

```

```

train_data_x = as.matrix(train_data[,-6])
train_data_y = as.matrix(train_data[, 6])
test_data_x = as.matrix(test_data[,-6])
test_data_y = as.matrix(test_data[, 6])

grid = 10^seq(10, -2, length = 100)

lasso.mod.cv = cv.glmnet(train_data_x, train_data_y, alpha = 1, lambda = grid, nfolds = 5)

bestlambda = lasso.mod.cv$lambda.min

lasso.fulldata = glmnet(train_data_x, train_data_y, alpha = 1, lambda = grid)
lasso.coef = predict(lasso.fulldata, type = "coefficients", s = bestlambda)

lasso.predict = predict(lasso.fulldata, s = bestlambda, newx = test_data_x)
mse.lasso = mean((test_data_y - lasso.predict)^2)
vectorMSElasso[j] = mse.lasso
}

mean(vectorMSElasso)

## [1] 0.1132495

```

Ridge

```

set.seed(1234)

vectorMSEridge = vector()
for (j in 1:7) {

  train_data = train[-folds[[j]], ]
  test_data = train[folds[[j]], ]

  train_data_x = as.matrix(train_data[,-6])
  train_data_y = as.matrix(train_data[, 6])
  test_data_x = as.matrix(test_data[,-6])
  test_data_y = as.matrix(test_data[, 6])

  grid = 10^seq(10, -2, length = 100)

  lasso.mod.cv = cv.glmnet(train_data_x, train_data_y, alpha = 0, lambda = grid, nfolds = 5)

  bestlambda = lasso.mod.cv$lambda.min

  lasso.fulldata = glmnet(train_data_x, train_data_y, alpha = 0, lambda = grid)
  lasso.coef = predict(lasso.fulldata, type = "coefficients", s = bestlambda)

  lasso.predict = predict(lasso.fulldata, s = bestlambda, newx = test_data_x)
  mse.lasso = mean((test_data_y - lasso.predict)^2)
  vectorMSEridge[j] = mse.lasso
}

mean(vectorMSEridge)

```

```
## [1] 0.110535
#lowest mean squared error was Ridge, so we will run this over entire training dataset
train_data_x = as.matrix(train[,-6])
train_data_y = as.matrix(train[, 6])
test_data_x = as.matrix(test[,-6])
test_data_y = as.matrix(test[, 6])

grid = 10^seq(10, -2, length = 100)

ridge.mod.cv = cv.glmnet(train_data_x, train_data_y, alpha = 0, lambda = grid, nfolds = 5)

bestlambda = ridge.mod.cv$lambda.min

lasso.fulldata = glmnet(train_data_x, train_data_y, alpha = 0, lambda = grid)
ridge.coef = predict(lasso.fulldata, type = "coefficients", s = bestlambda)

ridge.predict = predict(lasso.fulldata, s = bestlambda, newx = test_data_x)
mse.ridge = mean((test_data_y - ridge.predict)^2)
paste("mean squared error on test data for ridge is ", mse.ridge)

## [1] "mean squared error on test data for ridge is 0.0980150942822328"
paste("best lambda for ridge is ", bestlambda)

## [1] "best lambda for ridge is 0.01"
```

The Optimal model

Using our various models we produced the optimal model by selecting the parameters ridge identified as the most valuable. We use a number of thresholds for eliminating variables and identify the elbow at which the reduction in MSE becomes negligible.

```
thresholds = seq(1.0e-07, 5.0e-01, length.out = 100)
mse_values = c(length(thresholds))
num_non_zero_coefs = c(length(thresholds)) # Vector to store number of non-zero coefficients

# Loop over thresholds and calculate MSE for each using normal linear regression
for (i in 1:length(thresholds)) {
  # Extract coefficients, excluding the intercept
  ridge.coef = coef(lasso.fulldata, s = bestlambda)
  coef_thresholded = as.numeric(ridge.coef[-1]) # Remove intercept

  # Apply thresholding
  coef_thresholded[abs(coef_thresholded) < thresholds[i]] = 0

  # Identify selected predictors
  selected_predictors = which(coef_thresholded != 0)
  num_non_zero_coefs[i] = length(selected_predictors)

  if (length(selected_predictors) > 0) {
    # Subset training and test data matrices
```

```

reduced_train_data_x = train_data_x[, selected_predictors, drop = FALSE]
reduced_test_data_x = test_data_x[, selected_predictors, drop = FALSE]

# Convert to data frames for lm()
reduced_train_data_x_df = cbind(as.data.frame(reduced_train_data_x), loglatestPrice = train_data_y)
reduced_test_data_x_df = cbind(as.data.frame(reduced_test_data_x), loglatestPrice = test_data_y)

#print(names(reduced_train_data_x_df)) # Should include 'loglatestPrice'

#print("BREAK BREAK BREAK ")

#print(names(reduced_test_data_x_df)) # Should include 'loglatestPrice'

# Fit linear regression model
lm.reduced = lm(loglatestPrice ~ ., data = reduced_train_data_x_df)
lm.predict = predict(lm.reduced, newdata = reduced_test_data_x_df)

# Compute MSE
mse_values[i] = mean((reduced_test_data_x_df$loglatestPrice - lm.predict)^2)
} else {
mse_values[i] = NA # No predictors selected
}
}

#via analyzing the graph we get the following optimal threshold
opt.threshold = NULL
opt.threshold.num = NULL

# Loop through the MSE values and find the threshold just before MSE exceeds 0.2
for (i in 2:length(mse_values)) {
  if (mse_values[i] > 0.15 && mse_values[i - 1] <= 0.15) {
    opt_threshold = thresholds[i - 1]
    opt_threshold_num = i-1
    break
  }
}

# Print the threshold just before MSE exceeds 0.21
paste('We thus having that by finding the elbow joint, the optimal threshold is about', opt_threshold,

## [1] "We thus having that by finding the elbow joint, the optimal threshold is about 0.0959596767676767

# Adjust margins to create more space for the second y-axis label
par(mar = c(5, 4, 4, 6))

# Plot MSE and number of non-zero predictors on the same plot
plot(thresholds, mse_values, type = "l", col = "blue",
     xlab = "Threshold (Coefficient Magnitude)",
     ylab = "Mean Squared Error",
     main = "MSE and Number of Non-Zero Predictors vs. Threshold")

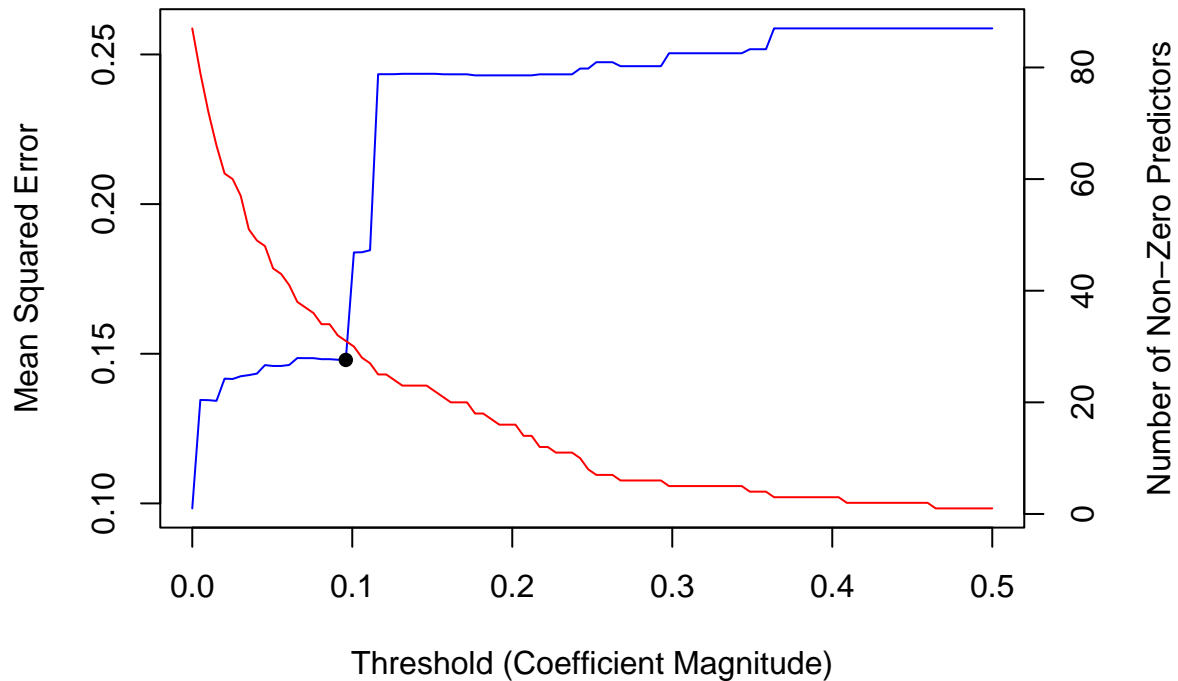
# Add a black dot at a specific index (e.g., opt_index)
points(thresholds[opt_threshold_num], mse_values[opt_threshold_num], col = "black", pch = 16) # pch =

```

```
# Add the second y-axis for the number of non-zero predictors
par(new = TRUE)
plot(thresholds, num_non_zero_coefs, type = "l", col = "red",
     axes = FALSE, xlab = "", ylab = "")

# Label the right axis
axis(side = 4)
mtext("Number of Non-Zero Predictors", side = 4, line = 3)
```

MSE and Number of Non-Zero Predictors vs. Threshold



```
print(thresholds[opt_threshold_num])
```

```
## [1] 0.09595968
```

```
print(mse_values[opt_threshold_num])
```

```
## [1] 0.1479377
```

```
# Reset plot parameters
par(mfrow = c(1, 1))
```

Analyzing the results from the vector

Comparison with Naive Model

As noted in our report we now consider a model based on what some relevant field experts consider to be the main features most buyers use when valuing houses.

```
# Fit the normal linear regression model for naive data
lm_model_naive = lm(train_naive$loglatestPrice ~ ., data = train_naive)

naive_predictions <- predict(lm_model_naive, newdata = test_naive)
```

```
mse_naive <- mean((test_naive$loglatestPrice - naive_predictions)^2)

paste('Our coefficients form a linear model with MSE', mse_values[opt_threshold_num], 'and the naive model produces MSE', mse_naive)

## [1] "Our coefficients form a linear model with MSE 0.147937667263368 and the naive model produces MSE 0.152062332736632"
```

Residual Plots

```
# Residual plot for the naive model
par(mfrow = c(1, 2)) # Set up two plots side by side

plot(naive_predictions, test_naive$loglatestPrice - naive_predictions,
     main = "Residuals for Naive Model",
     xlab = "Predicted Values",
     ylab = "Residuals")
abline(h = 0, col = "red")

# Fit the optimal model
coef_thresholded = ridge.coef
coef_thresholded = as.numeric(ridge.coef[-1]) # Remove intercept

coef_thresholded[abs(coef_thresholded) < thresholds[opt_threshold_num]] = 0

opt_predictors = which(coef_thresholded != 0)

opt_reduced_train_data_x = train[, opt_predictors, drop = FALSE]
opt_reduced_test_data_x = test[, opt_predictors, drop = FALSE]

opt_reduced_train_data_x_df = cbind(as.data.frame(opt_reduced_train_data_x), loglatestPrice = train_data$loglatestPrice)
opt_reduced_test_data_x_df = cbind(as.data.frame(opt_reduced_test_data_x), loglatestPrice = test_data$loglatestPrice)

lm_model_optimal = lm(loglatestPrice ~ ., data = opt_reduced_train_data_x_df)
opt_predictions = predict(lm_model_optimal, newdata = opt_reduced_test_data_x_df)

paste("MSE: ", mean((opt_reduced_test_data_x_df$loglatestPrice - opt_predictions)^2))

## [1] "MSE: 0.12685782528971"

summary(lm_model_optimal)
```

```
##
## Call:
## lm(formula = loglatestPrice ~ ., data = opt_reduced_train_data_x_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.86104 -0.24378 -0.04368  0.20947  2.42344
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.206e+01  1.159e-02 1040.480 < 2e-16 ***
## avgSchoolRating  6.688e-02  2.047e-03  32.671 < 2e-16 ***
## logavgSchoolDistance -1.538e-01  6.520e-03 -23.596 < 2e-16 ***
## livingMult      2.251e-03  3.172e-05   70.951 < 2e-16 ***
```



```
## month12                5.212e-02  1.300e-02   4.008 6.15e-05 ***
## propertyTaxRate2.01    -9.879e-02  1.872e-02  -5.278 1.33e-07 ***
## parkingSpaces4         7.725e-02  1.765e-02   4.377 1.21e-05 ***
## parkingSpaces8         1.439e-01  1.551e-01   0.928  0.3535
## parkingSpaces10        1.819e-01  1.900e-01   0.957  0.3384
## numPriceChanges4       -1.759e-02  1.122e-02  -1.567  0.1171
## numPriceChanges5       -8.002e-02  1.367e-02  -5.854 4.93e-09 ***
## numPriceChanges6       -1.476e-01  1.626e-02  -9.073 < 2e-16 ***
## numPriceChanges7       -2.399e-01  1.976e-02 -12.140 < 2e-16 ***
## numPriceChanges8       -1.699e-01  2.431e-02  -6.991 2.86e-12 ***
## numPriceChanges9       -1.825e-01  3.076e-02  -5.933 3.06e-09 ***
## numPriceChanges11      -1.372e-02  4.762e-02  -0.288  0.7732
## numPriceChanges12      -4.992e-02  4.663e-02  -1.071  0.2843
## numPriceChanges13      -8.743e-02  5.611e-02  -1.558  0.1192
## numPriceChanges14      -3.060e-02  7.181e-02  -0.426  0.6701
## numPriceChanges17       5.368e-02  1.898e-01   0.283  0.7773
## numPriceChanges19       4.857e-02  2.190e-01   0.222  0.8245
## numOfAppliances9       -4.717e-02  5.184e-02  -0.910  0.3629
## numOfPatioAndPorchFeatures3 1.306e-01  1.578e-02   8.276 < 2e-16 ***
## numOfPatioAndPorchFeatures5 1.123e-01  8.496e-02   1.321  0.1864
## numOfPatioAndPorchFeatures6 3.433e-01  1.897e-01   1.809  0.0704 .
## numOfPatioAndPorchFeatures7 2.726e-02  2.694e-01   0.101  0.9194
## numOfSecurityFeatures4   7.460e-02  4.632e-02   1.611  0.1073
## numOfSecurityFeatures6  -5.671e-02  2.682e-01  -0.211  0.8325
## numOfWaterfrontFeatures1 5.690e-01  1.096e-01   5.191 2.12e-07 ***
## numOfStories2          -1.296e-01  8.129e-03 -15.945 < 2e-16 ***
## numOfStories3          3.606e-01  3.684e-02   9.789 < 2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.3791 on 12658 degrees of freedom
```

```
## Multiple R-squared:  0.4812, Adjusted R-squared:  0.4799
```

```
## F-statistic: 391.3 on 30 and 12658 DF,  p-value: < 2.2e-16
```

```
optimal_predictions = predict(lm_model_optimal, newdata = opt_reduced_test_data_x_df)
```

```
# Plot residuals for the optimal model
```

```
plot(optimal_predictions, test_data_y - optimal_predictions,
```


```
  main = "Residuals for Optimal Model",
```

```
  xlab = "Predicted Values",
```


```
  ylab = "Residuals")
```

```
abline(h = 0, col = "red")
```

Residuals for Naive Model



Residuals for Optimal Model



```
# Reset the plot layout
par(mfrow = c(1, 1))
```

Whole Regression Model

```
train_data_x_df = cbind(as.data.frame(train_data_x), loglatestPrice = train_data_y)
test_data_x_df = cbind(as.data.frame(test_data_x), loglatestPrice = test_data_y)

lm_model_whole = lm(loglatestPrice ~ ., data = train_data_x_df)

predictions = predict(lm_model_whole, newdata = test_data_x_df)

paste('MSE: ', mean((test_data_x_df$loglatestPrice - predictions)^2))
```

```
## [1] "MSE: 0.098338136534761"
```

```
summary(lm_model_whole)
```

```
##
## Call:
## lm(formula = loglatestPrice ~ ., data = train_data_x_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.91937 -0.18793 -0.01875  0.16828  2.29747
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.582e+01  4.224e-01  37.447  < 2e-16 ***
```

| | | | | | |
|-----------------------------|------------|-----------|---------|----------|-----|
| ## yearBuilt | -1.459e-03 | 1.993e-04 | -7.324 | 2.55e-13 | *** |
| ## avgSchoolRating | 1.045e-01 | 2.259e-03 | 46.262 | < 2e-16 | *** |
| ## gdp_3_month_lag | 2.308e-07 | 5.712e-08 | 4.041 | 5.35e-05 | *** |
| ## unemployment_1_month_lag | 8.677e-04 | 1.840e-03 | 0.472 | 0.637175 | |
| ## mortgage_1_month_lag | -4.561e-02 | 9.508e-03 | -4.797 | 1.63e-06 | *** |
| ## loglotSizeSqFt | 1.079e-01 | 6.782e-03 | 15.916 | < 2e-16 | *** |
| ## logavgSchoolDistance | -7.648e-02 | 6.192e-03 | -12.350 | < 2e-16 | *** |
| ## logavgSchoolSize | -3.199e-01 | 1.419e-02 | -22.544 | < 2e-16 | *** |
| ## livingMult | 2.162e-03 | 3.386e-05 | 63.834 | < 2e-16 | *** |
| ## hasAssociationTRUE | -1.793e-01 | 8.398e-03 | -21.351 | < 2e-16 | *** |
| ## hasCoolingTRUE | -3.795e-02 | 2.629e-02 | -1.444 | 0.148852 | |
| ## hasSpaTRUE | 5.807e-02 | 1.195e-02 | 4.858 | 1.20e-06 | *** |
| ## hasViewTRUE | 3.200e-02 | 7.536e-03 | 4.246 | 2.19e-05 | *** |
| ## month2 | 6.698e-03 | 1.858e-02 | 0.361 | 0.718437 | |
| ## month3 | 1.645e-02 | 1.772e-02 | 0.929 | 0.353165 | |
| ## month4 | 9.884e-03 | 1.750e-02 | 0.565 | 0.572315 | |
| ## month5 | -7.598e-03 | 1.760e-02 | -0.432 | 0.665998 | |
| ## month6 | -2.899e-03 | 1.748e-02 | -0.166 | 0.868248 | |
| ## month7 | -3.396e-04 | 1.735e-02 | -0.020 | 0.984385 | |
| ## month8 | 1.523e-02 | 1.748e-02 | 0.871 | 0.383602 | |
| ## month9 | 2.694e-02 | 1.776e-02 | 1.517 | 0.129343 | |
| ## month10 | -3.457e-03 | 1.772e-02 | -0.195 | 0.845321 | |
| ## month11 | 9.234e-03 | 1.812e-02 | 0.510 | 0.610408 | |
| ## month12 | 3.121e-02 | 1.811e-02 | 1.723 | 0.084859 | . |
| ## propertyTaxRate2.01 | -2.691e-01 | 1.856e-02 | -14.503 | < 2e-16 | *** |
| ## propertyTaxRate2.21 | -2.474e-01 | 1.281e-02 | -19.310 | < 2e-16 | *** |
| ## parkingSpaces1 | 6.183e-02 | 1.284e-02 | 4.816 | 1.49e-06 | *** |
| ## parkingSpaces2 | 1.110e-02 | 7.502e-03 | 1.480 | 0.138957 | |
| ## parkingSpaces3 | 2.196e-02 | 1.384e-02 | 1.587 | 0.112533 | |
| ## parkingSpaces4 | 1.816e-02 | 1.651e-02 | 1.100 | 0.271303 | |
| ## parkingSpaces5 | 1.047e-01 | 4.264e-02 | 2.454 | 0.014137 | * |
| ## parkingSpaces6 | -5.468e-02 | 3.838e-02 | -1.425 | 0.154236 | |
| ## parkingSpaces7 | 5.574e-02 | 1.182e-01 | 0.471 | 0.637300 | |
| ## parkingSpaces8 | 1.930e-02 | 1.355e-01 | 0.142 | 0.886754 | |
| ## parkingSpaces9 | -2.610e-01 | 1.915e-01 | -1.363 | 0.172900 | |
| ## parkingSpaces10 | 8.762e-02 | 1.661e-01 | 0.528 | 0.597821 | |
| ## parkingSpaces12 | 2.296e-01 | 1.933e-01 | 1.188 | 0.234880 | |
| ## numPriceChanges2 | 1.405e-02 | 8.102e-03 | 1.734 | 0.083031 | . |
| ## numPriceChanges3 | -4.138e-03 | 9.401e-03 | -0.440 | 0.659834 | |
| ## numPriceChanges4 | -3.905e-02 | 1.059e-02 | -3.686 | 0.000229 | *** |
| ## numPriceChanges5 | -1.117e-01 | 1.262e-02 | -8.855 | < 2e-16 | *** |
| ## numPriceChanges6 | -1.845e-01 | 1.477e-02 | -12.491 | < 2e-16 | *** |
| ## numPriceChanges7 | -2.777e-01 | 1.776e-02 | -15.639 | < 2e-16 | *** |
| ## numPriceChanges8 | -2.274e-01 | 2.166e-02 | -10.498 | < 2e-16 | *** |
| ## numPriceChanges9 | -2.216e-01 | 2.722e-02 | -8.140 | 4.34e-16 | *** |
| ## numPriceChanges10 | -1.965e-01 | 3.122e-02 | -6.293 | 3.21e-10 | *** |
| ## numPriceChanges11 | -7.408e-02 | 4.184e-02 | -1.771 | 0.076639 | . |
| ## numPriceChanges12 | -1.246e-01 | 4.103e-02 | -3.036 | 0.002405 | ** |
| ## numPriceChanges13 | -1.602e-01 | 4.924e-02 | -3.253 | 0.001144 | ** |
| ## numPriceChanges14 | -1.398e-01 | 6.340e-02 | -2.205 | 0.027483 | * |
| ## numPriceChanges15 | -1.688e-01 | 7.903e-02 | -2.137 | 0.032651 | * |
| ## numPriceChanges16 | 8.366e-02 | 1.105e-01 | 0.757 | 0.448990 | |
| ## numPriceChanges17 | -1.061e-02 | 1.660e-01 | -0.064 | 0.949045 | |
| ## numPriceChanges18 | -2.184e-01 | 2.350e-01 | -0.929 | 0.352789 | |

```

## numPriceChanges19      6.670e-02  1.912e-01  0.349 0.727166
## numPriceChanges20      2.064e-01  3.312e-01  0.623 0.533148
## numPriceChanges22      8.341e-02  3.336e-01  0.250 0.802586
## numPriceChanges23     -2.713e-02  3.311e-01 -0.082 0.934707
## numOfAppliances1       1.832e-02  1.680e-02  1.090 0.275572
## numOfAppliances2       4.516e-02  1.479e-02  3.055 0.002258 **
## numOfAppliances3       4.653e-02  1.379e-02  3.374 0.000744 ***
## numOfAppliances4       6.457e-02  1.399e-02  4.617 3.94e-06 ***
## numOfAppliances5       7.601e-02  1.629e-02  4.666 3.10e-06 ***
## numOfAppliances6       6.460e-02  1.876e-02  3.444 0.000574 ***
## numOfAppliances7       2.166e-02  1.837e-02  1.179 0.238224
## numOfAppliances8       5.943e-02  2.161e-02  2.751 0.005958 **
## numOfAppliances9      -2.570e-02  4.702e-02 -0.547 0.584663
## numOfAppliances10      2.065e-01  3.318e-01  0.622 0.533645
## numOfPatioAndPorchFeatures1 -1.102e-02  9.977e-03 -1.104 0.269592
## numOfPatioAndPorchFeatures2  3.890e-02  1.148e-02  3.387 0.000708 ***
## numOfPatioAndPorchFeatures3  5.007e-02  1.608e-02  3.114 0.001847 **
## numOfPatioAndPorchFeatures4  1.033e-01  2.918e-02  3.541 0.000400 ***
## numOfPatioAndPorchFeatures5  1.870e-02  7.528e-02  0.248 0.803782
## numOfPatioAndPorchFeatures6  2.456e-01  1.661e-01  1.478 0.139355
## numOfPatioAndPorchFeatures7 -1.059e-01  2.384e-01 -0.444 0.656966
## numOfPatioAndPorchFeatures8  4.640e-01  3.427e-01  1.354 0.175707
## numOfSecurityFeatures1     -1.421e-02  9.160e-03 -1.551 0.120876
## numOfSecurityFeatures2      1.556e-02  1.236e-02  1.259 0.208018
## numOfSecurityFeatures3      7.713e-03  1.883e-02  0.410 0.682084
## numOfSecurityFeatures4      7.617e-02  4.162e-02  1.830 0.067281 .
## numOfSecurityFeatures5      1.575e-01  8.645e-02  1.822 0.068519 .
## numOfSecurityFeatures6     -6.103e-02  2.348e-01 -0.260 0.794934
## numOfWaterfrontFeatures1    3.539e-01  9.598e-02  3.687 0.000227 ***
## numOfWaterfrontFeatures2    7.728e-01  1.358e-01  5.692 1.29e-08 ***
## numOfStories2            -2.675e-02  7.537e-03 -3.548 0.000389 ***
## numOfStories3            3.484e-01  3.255e-02 10.705 < 2e-16 ***
## numOfStories4            3.921e-01  2.368e-01  1.656 0.097780 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3305 on 12601 degrees of freedom
## Multiple R-squared:  0.6074, Adjusted R-squared:  0.6047
## F-statistic: 224.1 on 87 and 12601 DF, p-value: < 2.2e-16

# Plot residuals for the optimal model
plot(optimal_predictions, test_data_y - optimal_predictions,
     main = "Residuals for Whole Model",
     xlab = "Predicted Values",
     ylab = "Residuals")
abline(h = 0, col = "red")

```

Residuals for Whole Model

